

Question 1

$$\alpha_k = \frac{1}{k}$$

$$Q^1[s_{17}, right] \leftarrow Q^0[s_{17}, right] + \alpha_k((r + 0.9 \max_a Q^0[s_{18}, a']) - Q^0[s_{17}, right])$$

$$Q^1[s_{17}, right] \leftarrow 0 + \frac{1}{1} \times (2 + 0.9 \times 0 - 0) = 2$$

$$Q^1[s_{18}, up] \leftarrow Q^0[s_{18}, up] + \alpha_k((r + 0.9 \max_a Q^0[s_{14}, a']) - Q^0[s_{18}, up])$$

$$Q^1[s_{18}, up] \leftarrow 0 + \frac{1}{1} \times (8 + 0.9 \times 0 - 0) = 8$$

$$Q^1[s_{14}, right] \leftarrow Q^0[s_{14}, right] + \alpha_k((r + 0.9 \max_a Q^0[s_{15}, a']) - Q^0[s_{14}, right])$$

$$Q^1[s_{14}, right] \leftarrow 0 + \frac{1}{1} \times ((-6) + 0.9 \times 0 - 0) = -6$$

(b)

$$Q^1[s_{23}, up] \leftarrow Q^0[s_{23}, up] + \alpha_k((r + 0.9 \max_a Q^1[s_{18}, a']) - Q^0[s_{23}, up])$$

$$Q^1[s_{23}, up] \leftarrow 0 + \frac{1}{1} \times (0 + 0.9 \times 8 - 0) = 7.2$$

$$Q^2[s_{18}, up] \leftarrow Q^1[s_{18}, up] + \alpha_k((r + 0.9 \max_a Q^1[s_{14}, a']) - Q^1[s_{18}, up])$$

$$Q^2[s_{18}, up] \leftarrow 8 + \frac{1}{2} \times (0 + 0.9 \times 0 - 8) = 4$$

$$Q^2[s_{14}, right] \leftarrow Q^1[s_{14}, right] + \alpha_k((r + 0.9 \max_a Q^1[s_{15}, a']) - Q^1[s_{14}, right])$$

$$Q^2[s_{14}, right] \leftarrow -6 + \frac{1}{2} \times (10 + 0.9 \times 0 - (-6)) = 2$$

(c)

When running SARSA, $Q[s, a]$ is updated using $Q[s', a']$ instead of $\max_a Q[s', a']$. Therefore,

for the previous part, the difference is that when updating $Q^2[s_{18}, up]$, SARSA will use

$$Q^1[s_{14}, right] = -6 \text{ rather than } \max_a Q^1[s_{14}, a'] = 0.$$

Question 2

Part a

(1)

```
In [1]: import pandas as pd
import numpy as np
import math
from matplotlib import pyplot as plt
```

```
In [2]: rs_samples = pd.read_csv('samples/rs_1.csv', names= ['data'])
rs_samples.head()
```

```
Out[2]:
```


	data
0	-1
1	2
2	-1
3	-1
4	-1

```
In [3]: length = len(rs_samples)
```

```
In [4]: def approxProbability(samples):
    acceptedNum = 0
    trueRainNum = 0
    y = [None]
    answer = []
    for i in range(1,length):
        #1: accepted sample with R = T
        if samples[i] == 1:
            acceptedNum += 1
            trueRainNum += 1
        #2: accepted sample with R = F
        elif samples[i] == 2:
            acceptedNum += 1
        y.append(trueRainNum / acceptedNum)

    # answer[0] = y
    answer.append(y)
    # answer[1] = acceptedNum
    answer.append(acceptedNum)
    return answer
```

```
In [5]: x = rs_samples.index
y = approxProbability(list(rs_samples['data']))[0]
plt.plot(x,y,linewidth = 0.8)
plt.title('Approximation of  $P(r|s,w)$  based on the first N samples')
plt.xlabel('Number of generated samples N')
plt.ylabel('P(r|s,w)')
plt.xscale('log')
```



(2)

```
In [6]: approx = y[length - 1]
print("The algorithm's approximation of  $P(r|s,w)$  using 100000 samples is: {:.6}".format(approx))
```

The algorithm's approximation of $P(r|s,w)$ using 100000 samples is: 0.319814

Part b

(1)

Using Hoeffding's inequality to derive the tightest bound ϵ , we have:

$$\begin{aligned} 2e^{-2n\epsilon^2} &< \delta \\ e^{-2n\epsilon^2} &< \frac{\delta}{2} \\ \frac{1}{e^{2n\epsilon^2}} &< \frac{\delta}{2} \\ \frac{2}{\delta} &< e^{2n\epsilon^2} \\ \ln\left(\frac{2}{\delta}\right) &< \ln(e^{2n\epsilon^2}) \\ \ln\left(\frac{2}{\delta}\right) &< 2n\epsilon^2 \\ \epsilon^2 &> \frac{\ln\left(\frac{2}{\delta}\right)}{2n} \\ \epsilon &> \sqrt{\frac{\ln\left(\frac{2}{\delta}\right)}{2n}} \end{aligned}$$

(2)

```
In [7]: acceptedSampleNum = approxProbability(list(rs_samples['data']))[1]
print("At N=100000, there are n={} accepted samples.".format(acceptedSampleNum))
```

At N=100000, there are n=27910 accepted samples.

(3)

```
In [8]: def computeEpsilon(n):
        delta = 0.05
        return math.sqrt(math.log(2/delta)/(2*n))
```

Given $n = 27910$ and $\delta = 0.05$, we have:

```
In [9]: epsilon = computeEpsilon(27910)
print("The value of  $\epsilon$  for n=27910 is {:.6}.".format(epsilon))
```

The value of ϵ for n=27910 is 0.00812928.

(4)

```
In [10]: def upToPointEpsilons(samples):
          acceptedNum = 0
          e = [None]
          for i in range(1,length):
              if samples[i] == 1 or samples[i] == 2:
                  acceptedNum += 1
              currEpsilon = computeEpsilon(acceptedNum,)
              e.append(currEpsilon)

          return e
```

```
In [11]: def confidenceBounds(s, e):
    size = len(s)
    upperBound = [None]
    lowerBound = [None]
    bounds = []
    # probability should be in [0,1]
    for i in range(1,size):
        upperBound.append(min(s[i]+e[i],1)) # upper bound <=1
        lowerBound.append(max(s[i]-e[i],0)) # lower bound >=0

    # bounds[0] = upperBound
    bounds.append(upperBound)
    # bounds[1] = lowerBound
    bounds.append(lowerBound)

    return bounds
```

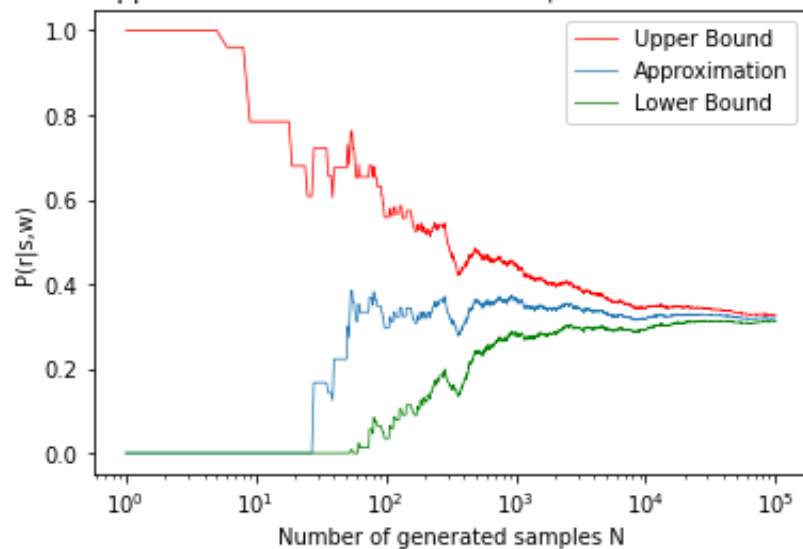
```
In [12]: s = approxProbability(list(rs_samples['data']))[0]
e = upToPointEpsilons(list(rs_samples['data']))
upperBound = confidenceBounds(s,e)[0]
lowerBound = confidenceBounds(s,e)[1]
```

```
In [13]: x = rs_samples.index

plt.plot(x,upperBound,linewidth = 0.7, c = 'r')
plt.plot(x,s,linewidth = 0.7)
plt.plot(x,lowerBound,linewidth = 0.7, c = 'g')

plt.legend(labels = ['Upper Bound', 'Approximation', 'Lower Bound'])
plt.title('Upper Bound, Approximation, Lower Bound of  $P(r|s,w)$  based on the first N samples')
plt.xlabel('Number of generated samples N')
plt.ylabel('P(r|s,w)')
plt.xscale('log')
```

Upper Bound, Approximation, Lower Bound of $P(r|s,w)$ based on the first N samples



Part c

(1)

```
In [14]: lw_samples = pd.read_csv('samples/lw_1.csv', names= ['samples', 'weights'])
# 1 based index
lw_samples.index += 1
lw_samples.head()
```

```
Out[14]:
```

	samples	weights
1	2	0.450
2	1	0.495
3	2	0.090
4	2	0.450
5	2	0.450


```
In [15]: # 1 denotes  $R=T$ 
# 2 denotes  $R=F$ 
def lwProbability(samples, weights):
    tSampleWeights = 0
    totalWeights = 0
    y = []
    for i in range(0, len(samples)):
        totalWeights += weights[i]
        if samples[i] == 1:
            tSampleWeights += weights[i]
        y.append(tSampleWeights/totalWeights)
    return y
```


```
In [16]: lw_y = lwProbability(list(lw_samples['samples']), list(lw_samples['weights']))
lwApprox = lw_y[len(lw_y) - 1]
print("The algorithm's approximation of  $P(r|s,w)$  using 100000 samples is: {:.6}".format(lwApprox))
```

The algorithm's approximation of $P(r|s,w)$ using 100000 samples is: 0.321788

(2)

```
In [17]: lw_x = lw_samples.index
lw_y = lwProbability(list(lw_samples['samples']), list(lw_samples['weights']))
plt.plot(lw_x, lw_y, linewidth = 0.7)
plt.title('Approximation of  $P(r|s,w)$  based on the first N samples using likelihood weighting')
plt.xlabel('Number of generated samples N')
plt.ylabel('P(r|s,w)')
plt.xscale('log')
```

Approximation of $P(r|s,w)$ based on the first N samples using likelihood weighting



(3)

The likelihood weighting algorithm seems to converge faster, especially when $N > 10^3$, whereas rejection sampling algorithm tends to converge when $N > 10^4$.

Question 3

(a)

We have the formula for forward message from filtering:

$$\begin{aligned} P(X_t|e_{0:t}) &= \alpha P(e_t|X_t)P(X_t|e_{0:t-1}) \\ &= \alpha P(e_t|X_t) \sum_{x_{t-1}} P(X_t|x_{t-1})P(x_{t-1}|e_{0:t-1}) \end{aligned}$$

Hence, when $t = 1$, we have:

$$\begin{aligned} P(X_1|e_{0:1}) &= P(X_1) = \sum_{x_0} P(X_1|x_0)P(x_0) \\ &= \langle 0.7, 0.3 \rangle \times 0.5 + \langle 0.4, 0.6 \rangle \times 0.5 \\ &= \langle 0.35, 0.15 \rangle + \langle 0.2, 0.3 \rangle \\ &= \langle 0.55, 0.45 \rangle \end{aligned}$$

$$\begin{aligned} P(X_1|e_1) &= \alpha P(e_1|X_1)P(X_1) \\ &= \alpha \langle 0.8, 0.3 \rangle \times \langle 0.55, 0.45 \rangle \\ &= \alpha \langle 0.44, 0.135 \rangle \\ &= \langle 0.7652173913, 0.2347826087 \rangle \end{aligned}$$

When $t = 2$, we have:

$$\begin{aligned} P(X_2|e_{0:1}) &= P(X_2|e_1) = \sum_{x_1} P(X_2|x_1)P(x_1|e_1) \\ &= \langle 0.7, 0.3 \rangle \times 0.7652173913 + \langle 0.4, 0.6 \rangle \times 0.2347826087 \\ &= \langle 0.535652173, 0.229565217 \rangle + \langle 0.093913043, 0.140869565 \rangle \\ &= \langle 0.629565216, 0.370434782 \rangle \end{aligned}$$

$$\begin{aligned} P(X_2|e_1, e_2) &= \alpha P(e_2|X_2)P(X_2|e_1) \\ &= \alpha \langle 0.2, 0.7 \rangle \times \langle 0.629565216, 0.370434782 \rangle \\ &= \alpha \langle 0.125913043, 0.259304347 \rangle \\ &= \langle 0.326862302, 0.673137697 \rangle \end{aligned}$$

When $t = 3$, we have

$$\begin{aligned} P(X_3|e_{0:2}) &= P(X_3|e_1, e_2) = \sum_{x_2} P(X_3|x_2)P(x_2|e_1, e_2) \\ &= \langle 0.7, 0.3 \rangle \times 0.326862302 + \langle 0.4, 0.6 \rangle \times 0.673137697 \\ &= \langle 0.228803611, 0.09805869 \rangle + \langle 0.269255078, 0.403882618 \rangle \\ &= \langle 0.498058689, 0.501941308 \rangle \end{aligned}$$

$$\begin{aligned} P(X_3|e_1, e_2, e_3) &= \alpha P(e_3|X_3)P(X_3|e_1, e_2) \\ &= \alpha \langle 0.8, 0.3 \rangle \times \langle 0.498058689, 0.501941308 \rangle \\ &= \alpha \langle 0.398446951, 0.150582392 \rangle \\ &= \langle 0.725729792, 0.274270207 \rangle \end{aligned}$$

For the backward message, we have the formula:

$$P(e_{k+1:t}|X_k) = \sum_{x_{k+1}} P(e_{k+1}|x_{k+1})P(e_{k+2:t}|x_{k+1})P(x_{k+1}|X_k)$$

Combining forward message and backward message, we have:

$$P(X_k|e_{0:t}) = \alpha P(X_k|e_{0:k})P(e_{k+1:t}|X_k)$$

Hence, when $k = 2$, we have:

$$\begin{aligned} P(e_3|X_2) &= \sum_{x_3} P(e_3|x_3)P(x_3|X_2) \\ &= 0.8 \times 1 \times \langle 0.7, 0.4 \rangle + 0.3 \times 1 \times \langle 0.3, 0.6 \rangle \\ &= \langle 0.56, 0.32 \rangle + \langle 0.09, 0.18 \rangle \\ &= \langle 0.65, 0.5 \rangle \end{aligned}$$

$$\begin{aligned} P(X_2|e_1, e_2, e_3) &= \alpha P(X_2|e_1, e_2)P(e_3|X_2) \\ &= \alpha \langle 0.326862302, 0.673137697 \rangle \times \langle 0.65, 0.5 \rangle \\ &= \alpha \langle 0.212460496, 0.336568848 \rangle \\ &= \langle 0.386974755, 0.613025244 \rangle \end{aligned}$$

When $k = 1$, we have:

$$\begin{aligned} P(e_2, e_3|X_1) &= \sum_{x_2} P(e_2|x_2)P(e_3|x_2)P(x_2|X_1) \\ &= 0.2 \times 0.65 \times \langle 0.7, 0.4 \rangle + 0.7 \times 0.5 \times \langle 0.3, 0.6 \rangle \\ &= \langle 0.091, 0.052 \rangle + \langle 0.105, 0.21 \rangle \\ &= \langle 0.196, 0.262 \rangle \end{aligned}$$

$$\begin{aligned} P(X_1|e_1, e_2, e_3) &= \alpha P(X_1|e_1)P(e_2, e_3|X_1) \\ &= \alpha \langle 0.7652173913, 0.2347826087 \rangle \times \langle 0.196, 0.262 \rangle \\ &= \alpha \langle 0.149982608, 0.061513043 \rangle \\ &= \langle 0.709152208, 0.290847791 \rangle \end{aligned}$$

$$P(X_1 = t|e_1, e_2, e_3) = 0.709152208$$

$$P(X_2 = t|e_1, e_2, e_3) = 0.386974755$$

$$P(X_3 = t|e_1, e_2, e_3) = 0.725729792$$

Hence, $P(X_i = t|\mathbf{e})$ for $i=1,2,3$ is $[0.709152208, 0.386974755, 0.725729792]$.

(b)

We have:

$$\max_{x_1, \dots, x_t} P(x_1, \dots, x_t, X_{t+1}, e_{1:t+1}) = P(e_{t+1}|x_{t+1}) \max_{x_t} [(P(X_{t+1}|x_t) m_{1:t})]$$

where

$$m_{1:t} = \max_{x_1, \dots, x_{t-1}} P(x_1, \dots, x_{t-1}, X_t, e_{1:t})$$

Using the result of filtering, when $t = 1$, we have:

$$m_{1:1} = P(X_1|e_1) = \langle 0.7652173913, 0.2347826087 \rangle$$

When $t=2$, we have:

$$\begin{aligned} m_{1:2} &= P(e_2|X_2) \langle \max [P(x_2|x_1) \times 0.7652173913, P(x_2|\sim x_1) \times 0.2347826087], \\ &\quad \max [P(\sim x_2|x_1) \times 0.7652173913, P(\sim x_2|\sim x_1) \times 0.2347826087] \rangle \\ &= \langle 0.2, 0.7 \rangle \times \langle \max [0.7 \times 0.7652173913, 0.4 \times 0.2347826087], \\ &\quad \max [0.3 \times 0.7652173913, 0.6 \times 0.2347826087] \rangle \\ &= \langle 0.2, 0.7 \rangle \times \langle \max(0.535652173, 0.093913043), \\ &\quad \max(0.229565217, 0.140871652) \rangle \\ &= \langle 0.2, 0.7 \rangle \times \langle 0.535652173, 0.229565217 \rangle \\ &= \langle 0.107130434, 0.160695651 \rangle \end{aligned}$$

When $t=3$, we have:

$$\begin{aligned} m_{1:3} &= P(e_3|X_3) \langle \max [P(x_3|x_2) \times 0.107130434, P(x_3|\sim x_2) \times 0.160695651], \\ &\quad \max [P(\sim x_3|x_2) \times 0.107130434, P(\sim x_3|\sim x_2) \times 0.160695651] \rangle \\ &= \langle 0.8, 0.3 \rangle \times \langle \max [0.7 \times 0.107130434, 0.4 \times 0.160695651], \\ &\quad \max [0.3 \times 0.107130434, 0.6 \times 0.160695651] \rangle \\ &= \langle 0.8, 0.3 \rangle \times \langle \max(0.074991303, 0.06427826), \max(0.03213913, 0.09641739) \rangle \\ &= \langle 0.8, 0.3 \rangle \times \langle 0.074991303, 0.09641739 \rangle \\ &= \langle 0.059993042, 0.028925217 \rangle \end{aligned}$$

Since $P(X_3 = t) > P(X_3 = f)$, it is most likely that $X_3 = t$.

$P(X_3 = t)$ traces back to $P(x_3|x_2)$, which shows that $X_3 = t$ is more likely from $X_2 = t$.

$P(X_2 = t)$ traces back to $P(x_2|x_1)$, which indicates that $X_2 = t$ is more likely from $X_1 = t$.

Hence, the most likely sequence of states of variables X_1, X_2, X_3 is (t,t,t).